

Virtual Memory

adapted for CS367@GMU

Bryant and O'Hallaron, Computer Systems: A Programmer's Perspective, Third Edition

1


Today

- **Basic Concepts & Address Spaces**
- VM as a tool for caching
- VM as a tool for memory management
- VM as a tool for memory protection
- Address translation

Bryant and O'Hallaron, Computer Systems: A Programmer's Perspective, Third Edition

2

A System Using Physical Addressing

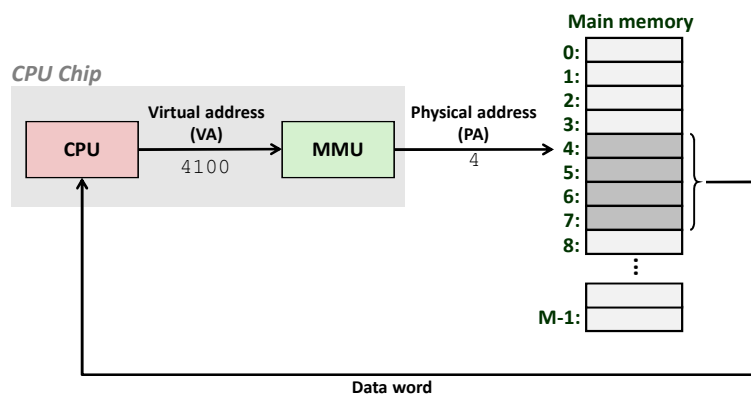


- Used in “simple” systems like embedded microcontrollers in devices like cars, elevators, and digital picture frames

Bryant and O'Hallaron, Computer Systems: A Programmer's Perspective, Third Edition

3

A System Using Virtual Addressing



- Memory Management Unit (MMU) performs the translation
- Used in all modern servers, laptops, and smart phones
- One of the great ideas in computer science

Bryant and O'Hallaron, Computer Systems: A Programmer's Perspective, Third Edition

4

Address Spaces

- **Linear address space:** Ordered set of contiguous non-negative integer addresses:
 $\{0, 1, 2, 3 \dots\}$
- **Virtual address space:** Set of $N = 2^n$ virtual addresses
 $V = \{0, 1, 2, 3, \dots, N-1\}$
- **Physical address space:** Set of $M = 2^m$ physical addresses
 $P = \{0, 1, 2, 3, \dots, M-1\}$
- **Clean distinction between data (bytes) and their attributes (addresses)**
- **Every byte in main memory has one physical address and zero or more virtual addresses**

Bryant and O'Hallaron, Computer Systems: A Programmer's Perspective, Third Edition

5

Address Spaces

- **Address Translation**
 - **MAP:** $V \rightarrow P \cup \{\emptyset\}$
 - For virtual address a :
 - **MAP(a) = a'** if data at virtual address a is at physical address a' in P
 - **MAP(a) = \emptyset** if data at virtual address a is invalid

Bryant and O'Hallaron, Computer Systems: A Programmer's Perspective, Third Edition

6

Why Virtual Memory (VM)?

- **Uses main memory efficiently**
 - Avoids ineffective use of DRAM (main memory) by reducing “fragmentation”
 - Uses DRAM as a cache for parts of a virtual address space
- **Simplifies memory management**
 - Each process gets the same uniform linear address space
 - Simplifies linking and loading
- **Isolates address spaces**
 - One process can't interfere with another's memory
 - User program cannot access privileged kernel information and code

Bryant and O'Hallaron, Computer Systems: A Programmer's Perspective, Third Edition

7

Memory at Run-time

- Each *process* (running program) has a private and distinct address space
- There are (typically) many different active processes sharing runtime memory.
- Need a way to enable this sharing in a way that is
 - Efficient – have to be able to quickly map from the process's address (virtual address) to the physical (real) address.
 - Safe – processes should not be able to interfere with each other.

Bryant and O'Hallaron, Computer Systems: A Programmer's Perspective, Third Edition


8

Contiguous Memory Allocation

A process can run if there is enough contiguous space in memory to hold its entire process space

Address Translation –
 $PA = VA + \text{offset}$ (where the process space starts in the memory)

Fragmentation – big problem!
 Process can't start to run because there is no large enough contiguous free region (consider Process D of size 800)





Bryant and O'Hallaron, Computer Systems: A Programmer's Perspective, Third Edition

9

Paged Memory

Divide physical address space in the main memory into equal-size chunks of size 2^n (called **Physical Pages**)

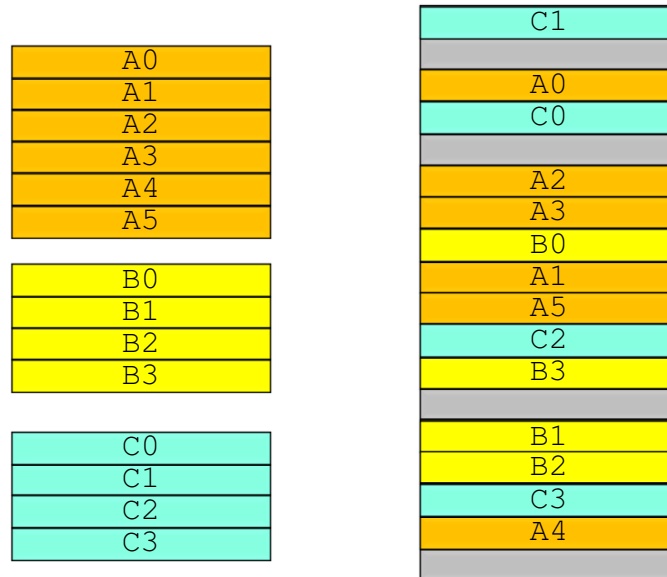
Divide the virtual address space of the process into equal-size chunks of size 2^n (called **Virtual Pages**)



Bryant and O'Hallaron, Computer Systems: A Programmer's Perspective, Third Edition

10

Paged Memory (Multiple Running Processes)



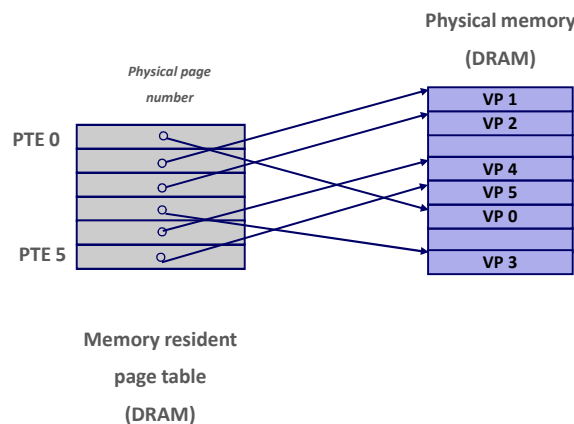
Bryant and O'Hallaron, Computer Systems: A Programmer's Perspective, Third Edition

11

Address Translation: Page Tables

A **page table** is an array of page table entries (PTEs) that maps virtual pages to physical pages.

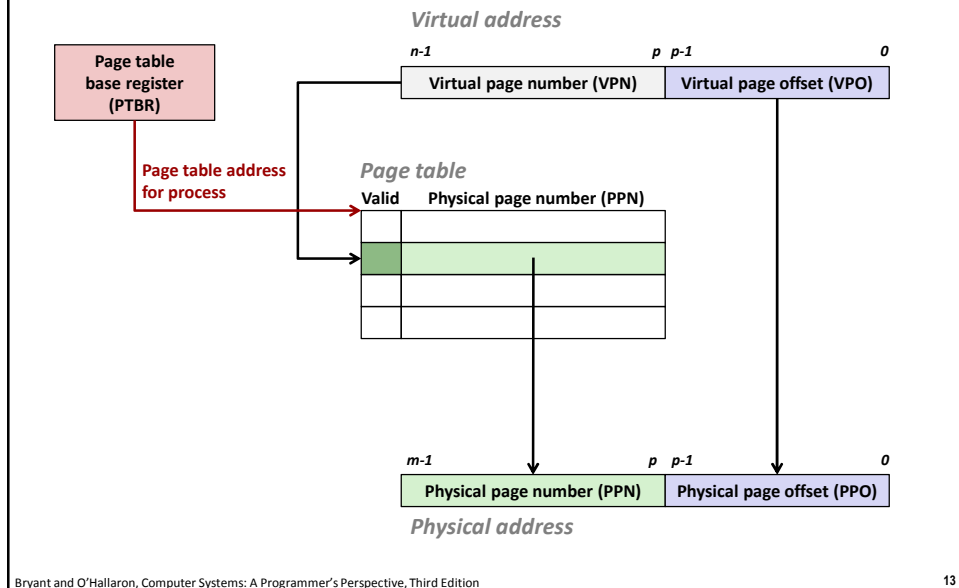
- Per-process data structure in DRAM (managed by the operating system)
- Dedicated register keeps address of page table in DRAM



Bryant and O'Hallaron, Computer Systems: A Programmer's Perspective, Third Edition

12

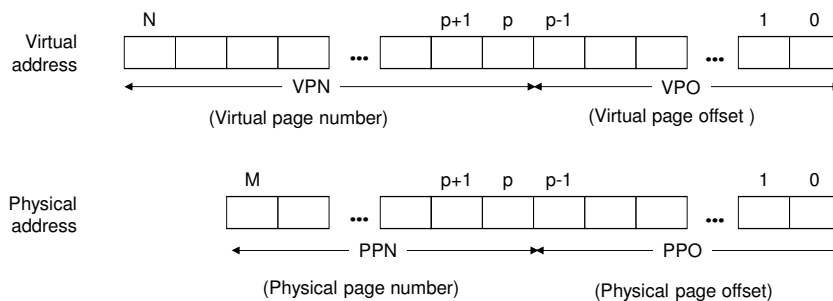
Address Translation With a Page Table



Address Translation

Size of VPN
- dependent on number of pages
for the process


Size of VPO = size of PPO
- dependent on physical page
size



Size of PPN
- dependent number of physical
pages in DRAM

Example:

**8-bit virtual address
with a 3-bit VPN ...**



What is the page size?


$$2^P = 2^5 = 32$$

How many pages can we have?

$$\text{Up to } 2^3 = 8$$

Bryant and O'Hallaron, Computer Systems: A Programmer's Perspective, Third Edition

15

Example:

Page Table:

000	0111100
001	1101100
010	0110101
011	1000100
100	0101110
101	0110000

How many bits in the physical address?

$$7_{ppn} + 5_{ppo} = 12 \text{ bits}$$

What is the physical address for:

• virtual address 01011011?

0110101 11011

• virtual address 10110001?

0110000 10001

6 entries <= 8 maximum: OK

Bryant and O'Hallaron, Computer Systems: A Programmer's Perspective, Third Edition

16

Today


- Address spaces & Basic Concepts
- VM as a tool for caching
- VM as a tool for memory management
- VM as a tool for memory protection
- Address translation

Bryant and O'Hallaron, Computer Systems: A Programmer's Perspective, Third Edition

17

VM as a Tool for Caching

- Conceptually, we can view the total *virtual memory* as an array of N contiguous bytes stored on disk at all times.
- The contents of the array on disk are cached in *physical memory* (which can be called the *DRAM cache*)
 - These cache blocks in this case correspond to *pages* (size is $P = 2^p$ bytes)



Bryant and O'Hallaron, Computer Systems: A Programmer's Perspective, Third Edition

18

DRAM Cache Organization

■ DRAM cache organization driven by the enormous miss penalty

- DRAM is about **10x** slower than SRAM
- Disk is about **10,000x** slower than DRAM

■ Consequences

- Large page (block) size: typically 4-8 KB, sometimes 4 MB
- Fully associative
 - Any VP can be placed in any PP
 - Requires a “large” mapping function – different from CPU caches
- Highly sophisticated, expensive replacement algorithms
 - Too complicated and open-ended to be implemented in hardware
- Write-back rather than write-through

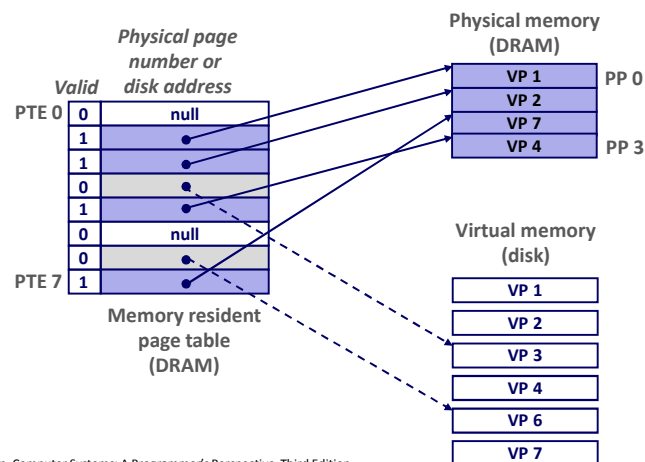
Bryant and O'Hallaron, Computer Systems: A Programmer's Perspective, Third Edition

19

Extending Page Table Structure

■ Page table entries (PTEs) may contain the disk address if VP is not in the physical memory (DRAM)

- Valid bit will be 0 for such entries

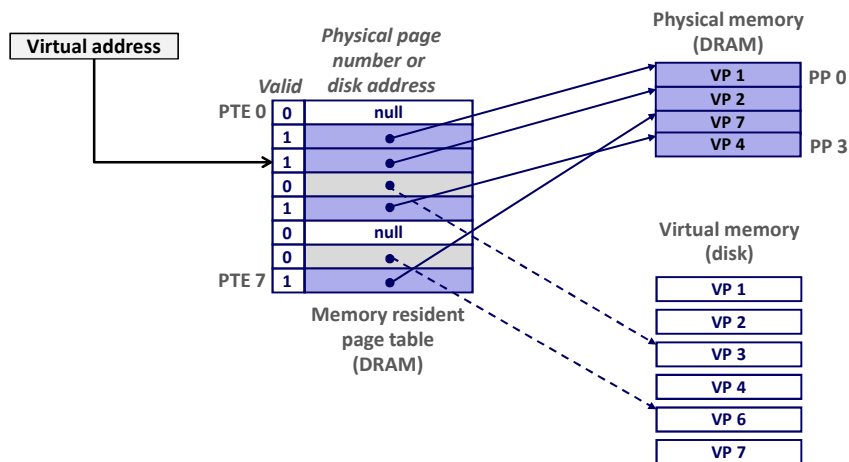


Bryant and O'Hallaron, Computer Systems: A Programmer's Perspective, Third Edition

20

Page Hit

- **Page hit:** reference to VM word that is in physical memory (DRAM cache hit)

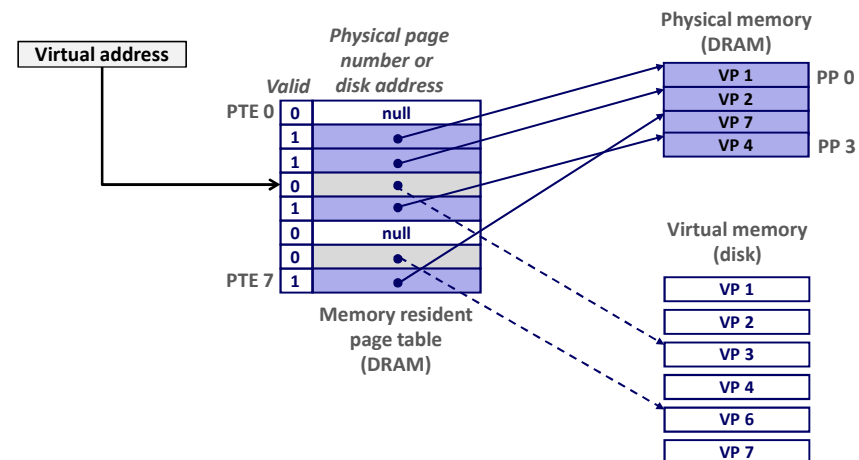


Bryant and O'Hallaron, Computer Systems: A Programmer's Perspective, Third Edition

21

Page Fault

- **Page fault:** reference to VM word that is not in physical memory (DRAM cache miss)

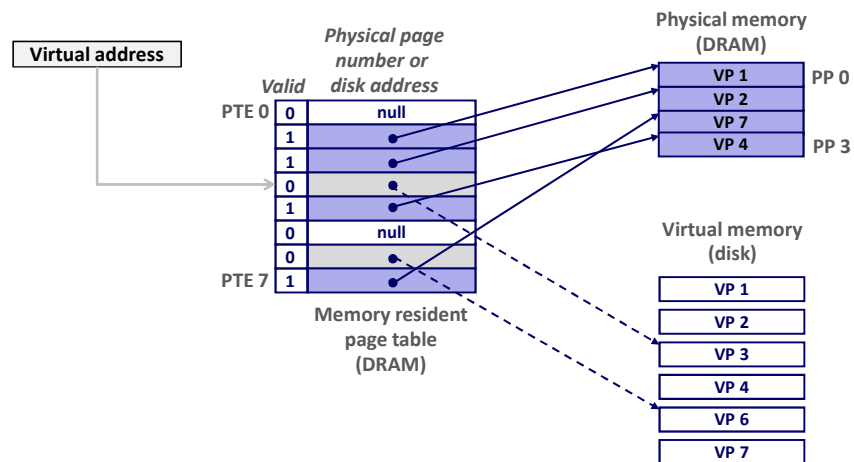


Bryant and O'Hallaron, Computer Systems: A Programmer's Perspective, Third Edition

22

Handling Page Fault

- Page miss causes **page fault** (an “exception” -- a transfer of control to the Operating System in response to some *event*)

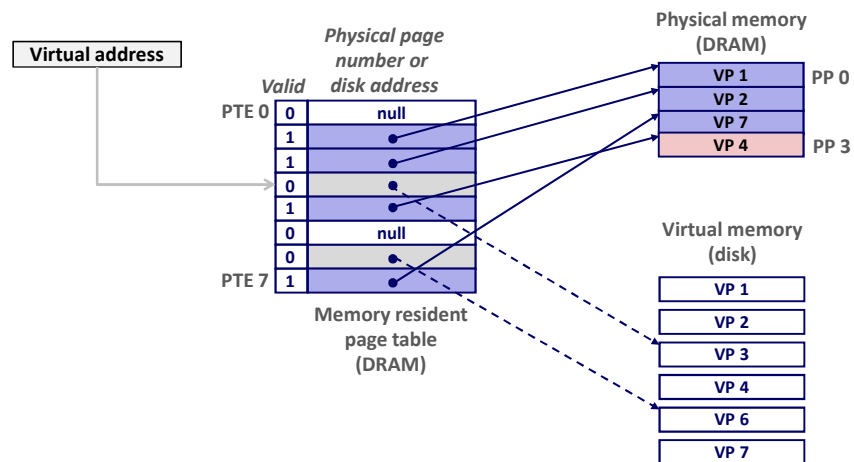


Bryant and O'Hallaron, Computer Systems: A Programmer's Perspective, Third Edition

23

Handling Page Fault

- Page miss causes page fault (an exception)
- Page fault handler selects a victim to be evicted (here VP 4)

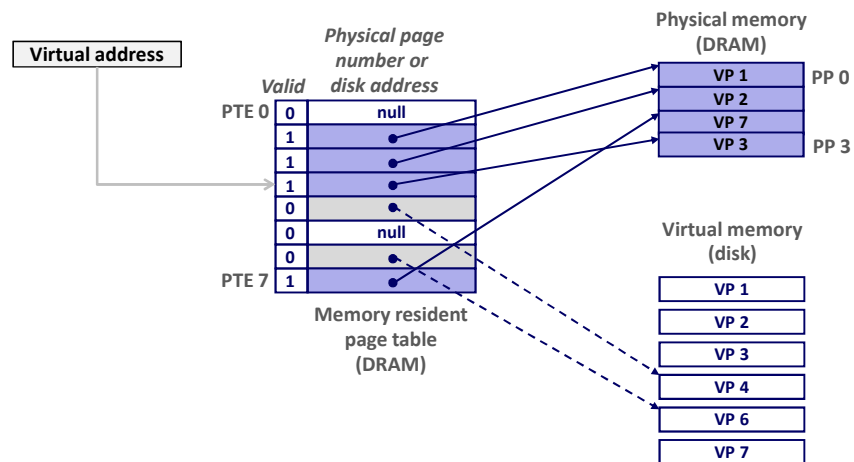


Bryant and O'Hallaron, Computer Systems: A Programmer's Perspective, Third Edition

24

Handling Page Fault

- Page miss causes page fault (an exception)
- Page fault handler selects a victim to be evicted (here VP 4)

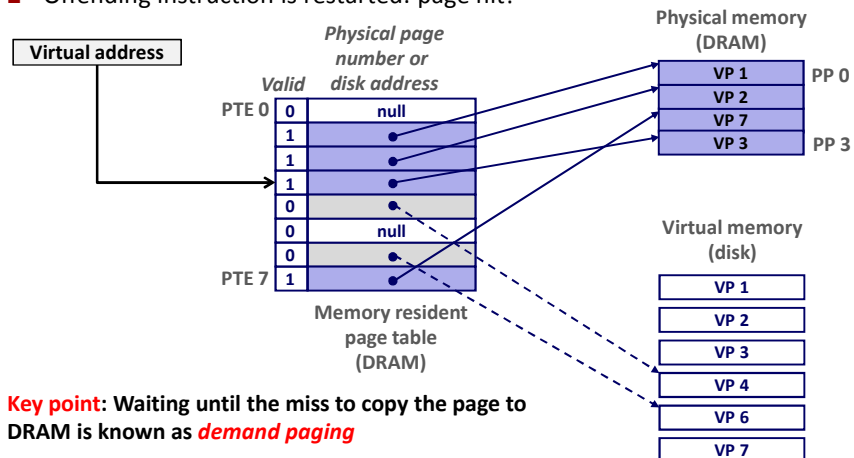


Bryant and O'Hallaron, Computer Systems: A Programmer's Perspective, Third Edition

25

Handling Page Fault

- Page miss causes page fault (an exception)
- Page fault handler selects a victim to be evicted (here VP 4)
- Offending instruction is restarted: page hit!



Key point: Waiting until the miss to copy the page to DRAM is known as *demand paging*

Bryant and O'Hallaron, Computer Systems: A Programmer's Perspective, Third Edition

26

Today


- Address spaces & Basic Concepts
- VM as a tool for caching
- **VM as a tool for memory management**
- VM as a tool for memory protection
- Address translation

Bryant and O'Hallaron, Computer Systems: A Programmer's Perspective, Third Edition

27

VM as a Tool for Memory Management

- **Key idea: each process has its own virtual address space**
 - It can view memory as a simple linear array
 - Mapping function scatters addresses through physical memory
 - Well-chosen mappings can improve locality



Bryant and O'Hallaron, Computer Systems: A Programmer's Perspective, Third Edition

28


VM as a Tool for Memory Management

■ Memory allocation

- Each virtual page can be mapped to any physical page
- A virtual page can be stored in different physical pages at different times

■ Sharing code and data among processes

- Map virtual pages to the same physical page (here: PP 6)



Bryant and O'Hallaron, Computer Systems: A Programmer's Perspective, Third Edition

29

Today

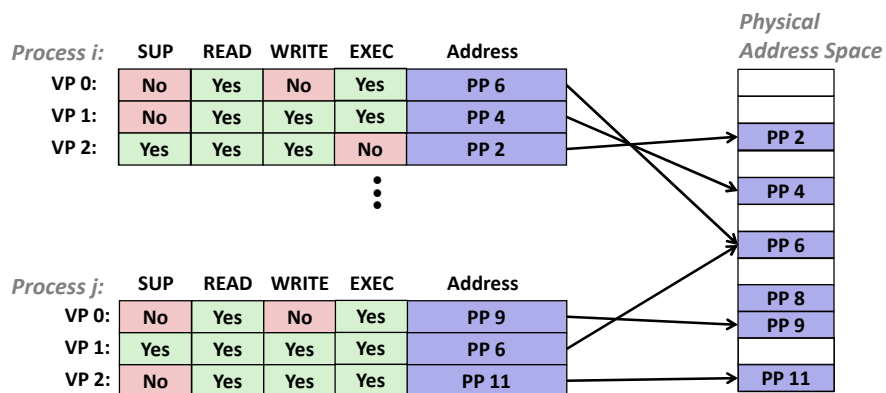
- Address spaces & Basic Concepts
- VM as a tool for caching
- VM as a tool for memory management
- VM as a tool for memory protection
- Address translation

Bryant and O'Hallaron, Computer Systems: A Programmer's Perspective, Third Edition

30

VM as a Tool for Memory Protection

- Extend PTEs with permission bits
- Page fault handler checks these before remapping
 - If violated, “segmentation fault” occurs!



Bryant and O'Hallaron, Computer Systems: A Programmer's Perspective, Third Edition

SUP: supervisor (kernel) mode

31

Today

- Address spaces & Basic Concepts
- VM as a tool for caching
- VM as a tool for memory management
- VM as a tool for memory protection
- Address translation

Bryant and O'Hallaron, Computer Systems: A Programmer's Perspective, Third Edition

32

VM Address Translation - Extended

- **Virtual Address Space**

- $V = \{0, 1, \dots, N-1\}$

- **Physical Address Space**

- $P = \{0, 1, \dots, M-1\}$

- **Address Translation**

- $MAP: V \rightarrow P \cup \{\emptyset\}$
 - For virtual address a :
 - $MAP(a) = a'$ if data at virtual address a is at physical address a' in P
 - $MAP(a) = \emptyset$ if data at virtual address a is not in physical memory
 - Either invalid or stored on disk

Bryant and O'Hallaron, Computer Systems: A Programmer's Perspective, Third Edition

33

Summary of Address Translation Symbols

- **Basic Parameters**

- $N = 2^n$: Number of addresses in virtual address space
 - $M = 2^m$: Number of addresses in physical address space
 - $P = 2^p$: Page size (bytes)

- **Components of the virtual address (VA)**

- TLBI: TLB index
 - TLBT: TLB tag
 - VPO: Virtual page offset
 - VPN: Virtual page number

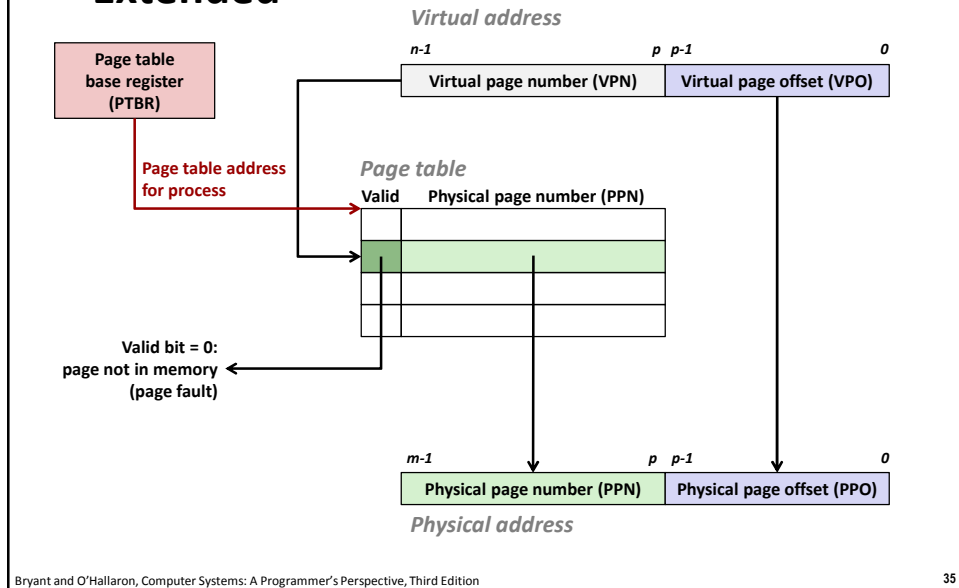
- **Components of the physical address (PA)**

- PPO: Physical page offset (same as VPO)
 - PPN: Physical page number

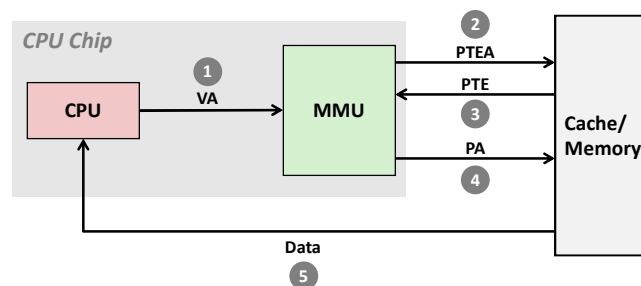
Bryant and O'Hallaron, Computer Systems: A Programmer's Perspective, Third Edition

34

Address Translation With a Page Table - Extended

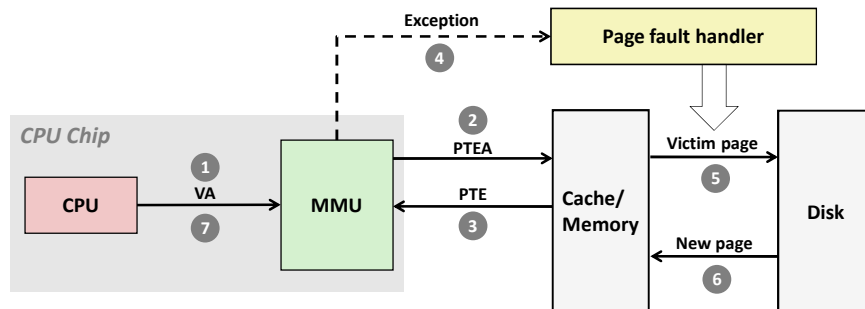


Address Translation: Page Hit



- 1) Processor sends virtual address to MMU
- 2-3) MMU fetches PTE from page table in memory
- 4) MMU sends physical address to cache/memory
- 5) Cache/memory sends data word to processor

Address Translation: Page Fault

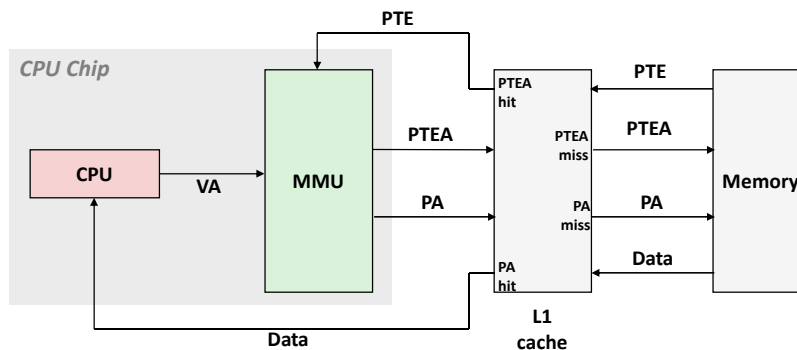


- 1) Processor sends virtual address to MMU
- 2-3) MMU fetches PTE from page table in memory
- 4) Valid bit is zero, so MMU triggers page fault exception
- 5) Handler identifies victim (and, if dirty, pages it out to disk)
- 6) Handler pages in new page and updates PTE in memory
- 7) Handler returns to original process, restarting faulting instruction

Bryant and O'Hallaron, Computer Systems: A Programmer's Perspective, Third Edition

37

Integrating VM and Cache



VA: virtual address, PA: physical address, PTE: page table entry, PTEA = PTE address

Bryant and O'Hallaron, Computer Systems: A Programmer's Perspective, Third Edition

38

Speeding up Translation with a TLB

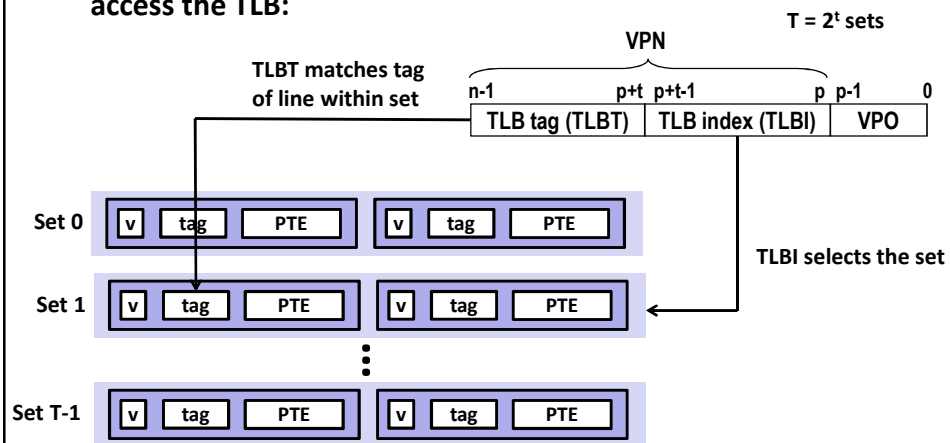
- Page table entries (PTEs) might be cached in L1 like any other memory word
 - PTEs may be evicted by other data references
 - PTE hit would still require a small L1 delay
- Solution: **Translation Lookaside Buffer (TLB)**
 - Small set-associative hardware cache in MMU
 - Maps virtual page numbers to physical page numbers
 - Contains complete page table entries for small number of pages

Bryant and O'Hallaron, Computer Systems: A Programmer's Perspective, Third Edition

39

Accessing the TLB

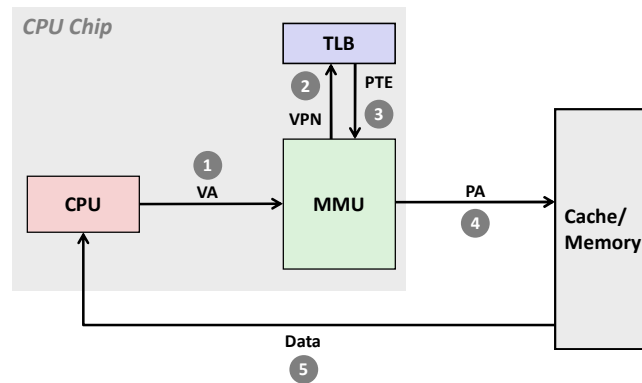
- MMU uses the VPN portion of the virtual address to access the TLB:



Bryant and O'Hallaron, Computer Systems: A Programmer's Perspective, Third Edition

40

TLB Hit

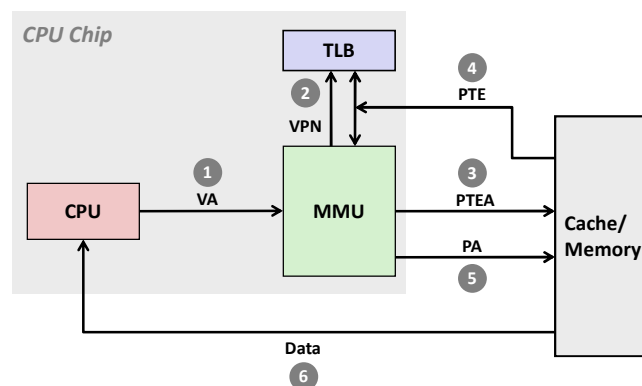


A TLB hit eliminates a memory access

Bryant and O'Hallaron, Computer Systems: A Programmer's Perspective, Third Edition

41

TLB Miss



A TLB miss incurs an additional memory access (the PTE)

Fortunately, TLB misses are rare: pages are large.

Bryant and O'Hallaron, Computer Systems: A Programmer's Perspective, Third Edition

42

Multi-Level Page Tables

- **Suppose:**

- 4KB (2^{12}) page size, 48-bit address space, 8-byte PTE


- **Problem:**

- Would need a 512 GB page table!
 - $2^{48} * 2^{-12} * 2^3 = 2^{39}$ bytes

- **Common solution: Multi-level page table**

- **Example: 2-level page table**

- Level 1 table: each PTE points to a page table (always memory resident)
 - Level 2 table: each PTE points to a page (paged in and out like any other data)



Bryant and O'Hallaron, Computer Systems: A Programmer's Perspective, Third Edition

43

Summary

- **Programmer's view of virtual memory**

- Each process has its own private linear address space
 - Cannot be corrupted by other processes

- **System view of virtual memory**

- Uses memory efficiently by "caching" virtual memory pages
 - Simplifies memory management and programming
 - Simplifies protection by providing a convenient interpositioning point to check permissions

Bryant and O'Hallaron, Computer Systems: A Programmer's Perspective, Third Edition

44